

# Automatic Hyper-parameter Tuning via Bayesian Optimization

**Janardhan Rao (Jana) Doppa**

School of EECS, Washington State University

# Motivation

- **Every machine learning algorithm has hyper-parameters**

- ▲ **Perceptron algorithm:** learning rate, number of iterations
- ▲ **SVM algorithm:** tradeoff parameter  $C$ , kernel parameter (degree for polynomial kernel, width for Gaussian kernel)
- ▲ **Random Forests:** number of trees, depth of each tree
- ▲ **Deep neural networks:** number of layers, weight regularization, layer size, which non-linearity, batch size, learning rate schedule, stopping conditions etc.
- ▲ ...

# Search for Good Hyper-parameters: Standard Practice

- **Define an objective function**
  - ▲ We care about generalization performance. Use cross-validation to measure parameter quality.
- **How do people currently search? Black magic.**
  - ▲ Grid search
  - ▲ Random search
  - ▲ Graduate student descent
- **What is wrong?**
  - ▲ Painful!
  - ▲ Computationally expensive – many training cycles
  - ▲ Possibly sub-optimal

# Bayesian Optimization: Key Ideas

- **Build a surrogate statistical model** based on past computational experiments
  - ▲ Assumption is that **surrogate model is cheap to evaluate**
- **Intelligently select the next experiment** (candidate solution) **using the statistical model**
  - ▲ **Trade-off exploration and exploitation**
  - ▲ Exploration corresponds to selecting candidates for which the statistical model is not confident (high variance)
  - ▲ Exploitation corresponds to selecting candidates for which the statistical model is highly confident (high mean)


# Bayesian Optimization Framework: Key Elements

- **Surrogate statistical model**
  - ▲ Cheap to evaluate
  - ▲ Can quantify uncertainty of predictions (i.e., variance)
- **Acquisition function**
  - ▲ Scores candidate solutions (via mean and variance obtained from the statistical model) in terms of their usefulness
- **Optimizer**
  - ▲ Select the candidate that maximizes the acquisition function
  - ▲ Next candidate we should try


# Bayesian Optimization Framework: High-level Overview

- Initialize statistical model  $F$
- Repeat the following steps for several iterations
  - ▲ Select the next candidate (say  $x$ ) by optimizing the acquisition function  $A(x)$
  - ▲ Run experiment with candidate  $x$  to compute its quality  $y$
  - ▲ Update the statistical model  $F$  based on the new training example  $(x, y)$
  - ▲ Update the best uncovered solution so far (say  $x_{\{best\}}$ )

# Bayesian Optimization: Illustration




# Bayesian Optimization: Illustration







# Bayesian Optimization: Illustration




# Bayesian Optimization: Illustration




# Bayesian Optimization: Illustration




# Bayesian Optimization: Illustration



# Bayesian Optimization: Illustration



# Bayesian Optimization: Illustration



# Bayesian Optimization Framework: Key Elements

- **Surrogate statistical model**
  - ▲ Cheap to evaluate
  - ▲ Can quantify uncertainty of predictions (i.e., variance)
- **Acquisition function**
  - ▲ Scores candidate solutions (via mean and variance obtained from the statistical model) in terms of their usefulness
- **Optimizer**
  - ▲ Select the candidate that maximizes the acquisition function
  - ▲ Next candidate we should try

# Surrogate Model: Choices

- **Bayesian Regression Model**
- **Gaussian Process (GP) Model**
  - ▲ Very popular
- **Random Forests**
  - ▲ A bunch of regression trees
  - ▲ Mean = empirical mean of the predictions from all trees
  - ▲ Variance = empirical variance of the predictions from all trees
  - ▲ Simple implement
  - ▲ Often works well. You can try Mondrian Forests (reason about uncertainty) for a better choice
  - ▲ Becoming very popular lately (see SMAC software)



# Acquisition Function: Choices


- Expected Improvement (EI)
- Probability of Improvement
- Upper Confidence Bound (UCB)
  
- EI and UCB are the most popular acquisition functions

# Optimizer: Choices

- Gradient descent with random restarts
- LBFGS
- Divided RECTangles (DIRECT) algorithm
- Simultaneous Optimistic Optimization (SOO) algorithm
- ...

# Example Results

CIFAR10: Deep convolutional neural net (Krizhevsky)  
Achieves 9.5% test error vs. 11% with hand tuning.



Snoek, Larochelle & RPA, NIPS 2012

# Papers and Software

- Jasper Snoek, Hugo Larochelle, Ryan P. Adams:  
Practical Bayesian Optimization of Machine Learning Algorithms. NIPS 2012: 2960-2968
  - ▲ <https://papers.nips.cc/paper/4522-practical-bayesian-optimization-of-machine-learning-algorithms.pdf>
- Spearmint Software
  - ▲ <https://github.com/JasperSnoek/spearmint>