

Lecture #4: LaSO Framework

Beam Search Framework

- **Given**

- ▶ Search space definition (ordered or unordered)
- ▶ Training examples (input-output pairs)
- ▶ Beam width B (>1)

- **Learning Goal**


- ▶ Learn a heuristic function to quickly guide the search to the correct “complete” output

- **Key Idea:**

- ▶ Structured prediction as a search problem in the space of partial outputs
- ▶ Training examples define target paths from initial state to the goal state (correct structured output)


Beam Search Framework: Key Elements

- 1) Search space; 2) Search procedure; 3) **Heuristic function**




- Represent heuristic function as a linear function
 - ▲ $H(n) = w \cdot \psi(n)$, where $\psi(n)$ stands for features of node n


Beam Search: Illustration




Beam Search: Illustration




Beam Search: Illustration



Beam Search: Illustration



Beam Search: Illustration



Beam Search Framework: Inference

- **Input:** learned weights w ; beam width B ; structured input x
- **repeat**
 - ▶ Perform search with heuristic $H(n) = w \cdot \psi(n)$
- **until** *reaching a terminal state*
- **Output:** the complete output y corresponding to the terminal state

Beam Search Framework: Generic Learning Template

- **Three design choices**

- ▲ How to define the notion of “search error”?
- ▲ How to “update the weights” of heuristic function when a search error is encountered?
- ▲ How to “update the beam” after weight update?

Beam Search Framework: Learning Instantiations

- Early update

[Collins and Roark, 2004]

- Max-violation update

[Huang et al., 2012]

- Learning as Search Optimization (LaSO)

[Daume et al., 2005], [Xu et al., 2009]

Beam Search Framework: Learning Instantiations

- Early update
- Max-violation update
- Learning as Search Optimization (LaSO)

Beam Search Framework: Early Update

- **Search error:** NO target node in the beam
 - ▲ We cannot reach the goal node (correct structured output)
- **Weight update:** standard structured perceptron
 - ▲ Score of correct output > score of bad output
- **Beam update:** reset beam with initial state OR discontinue search

Beam Search Framework: Early Update

- **repeat**
 - ▲ For every training example (x, y)
 - Perform search with current heuristic (weights)
 - If **search error**, **update weights**
 - Reset beam with initial state
 - (Dis)continue search
- **until** *convergence or max. iterations*

Beam Search Framework: Learning Instantiations

- Early update
- Max-violation update
- Learning as Search Optimization (LaSO)

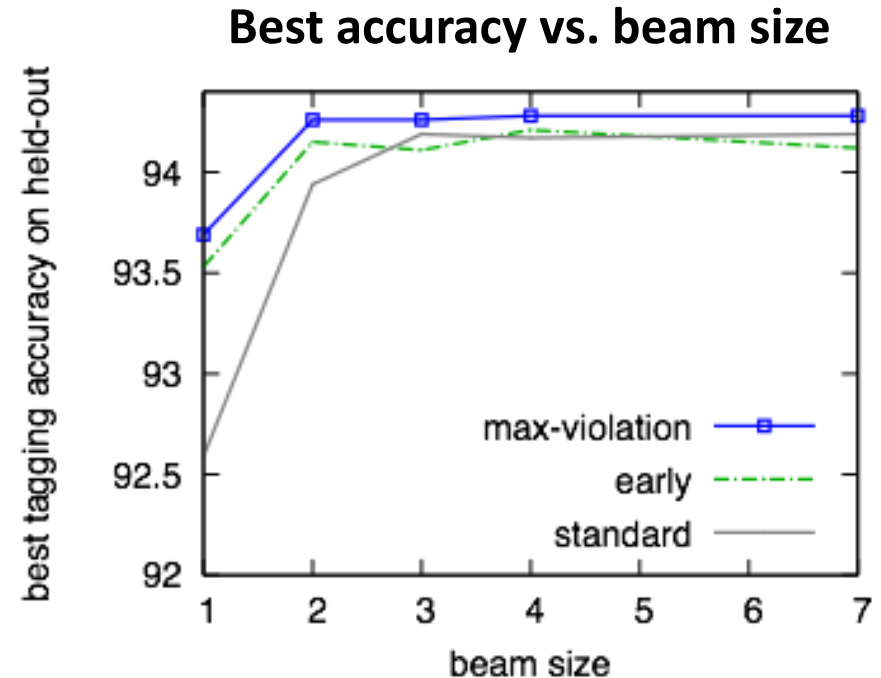
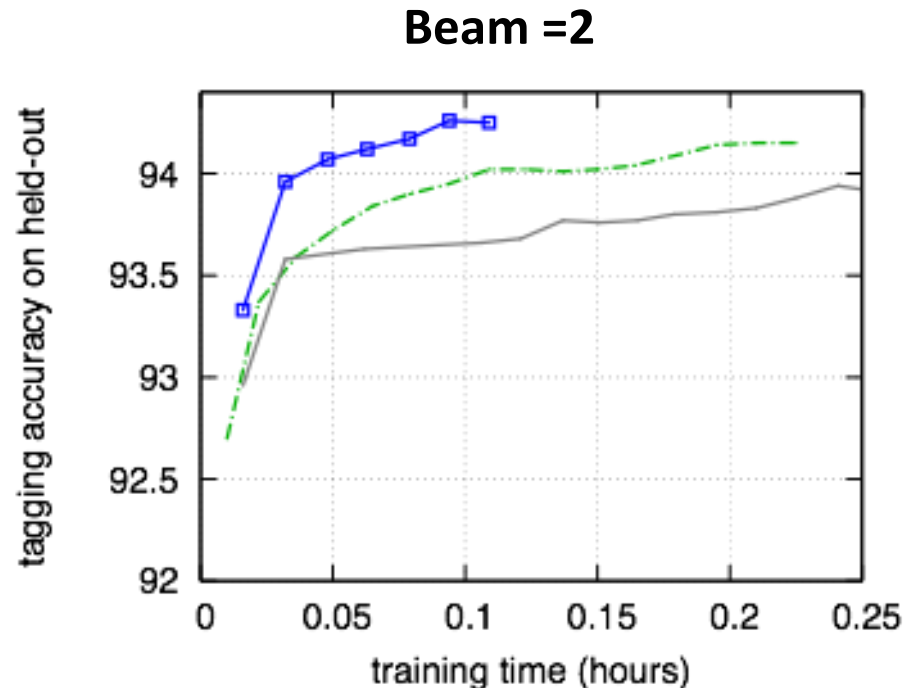
Beam Search Framework: Max-Violation Update

- Improves on the drawback of Early update
 - ▲ **Slow learning:** learns from only earliest mistake
- **Max-Violation fix**
 - ▲ Consider worst-mistake (maximum violation) instead of earliest-mistake for the weight update
 - ▲ More useful training data
 - ▲ Converges faster than early update

POS Tagging: Max-violation vs. Early vs. Standard

- **Early and Max-violation >> Standard at small beams**
 - ▲ Advantage shrinks as beam size increases
 - ▲ Max-violation converges faster than Early (and slightly better)

Source: Huang et al., 2012

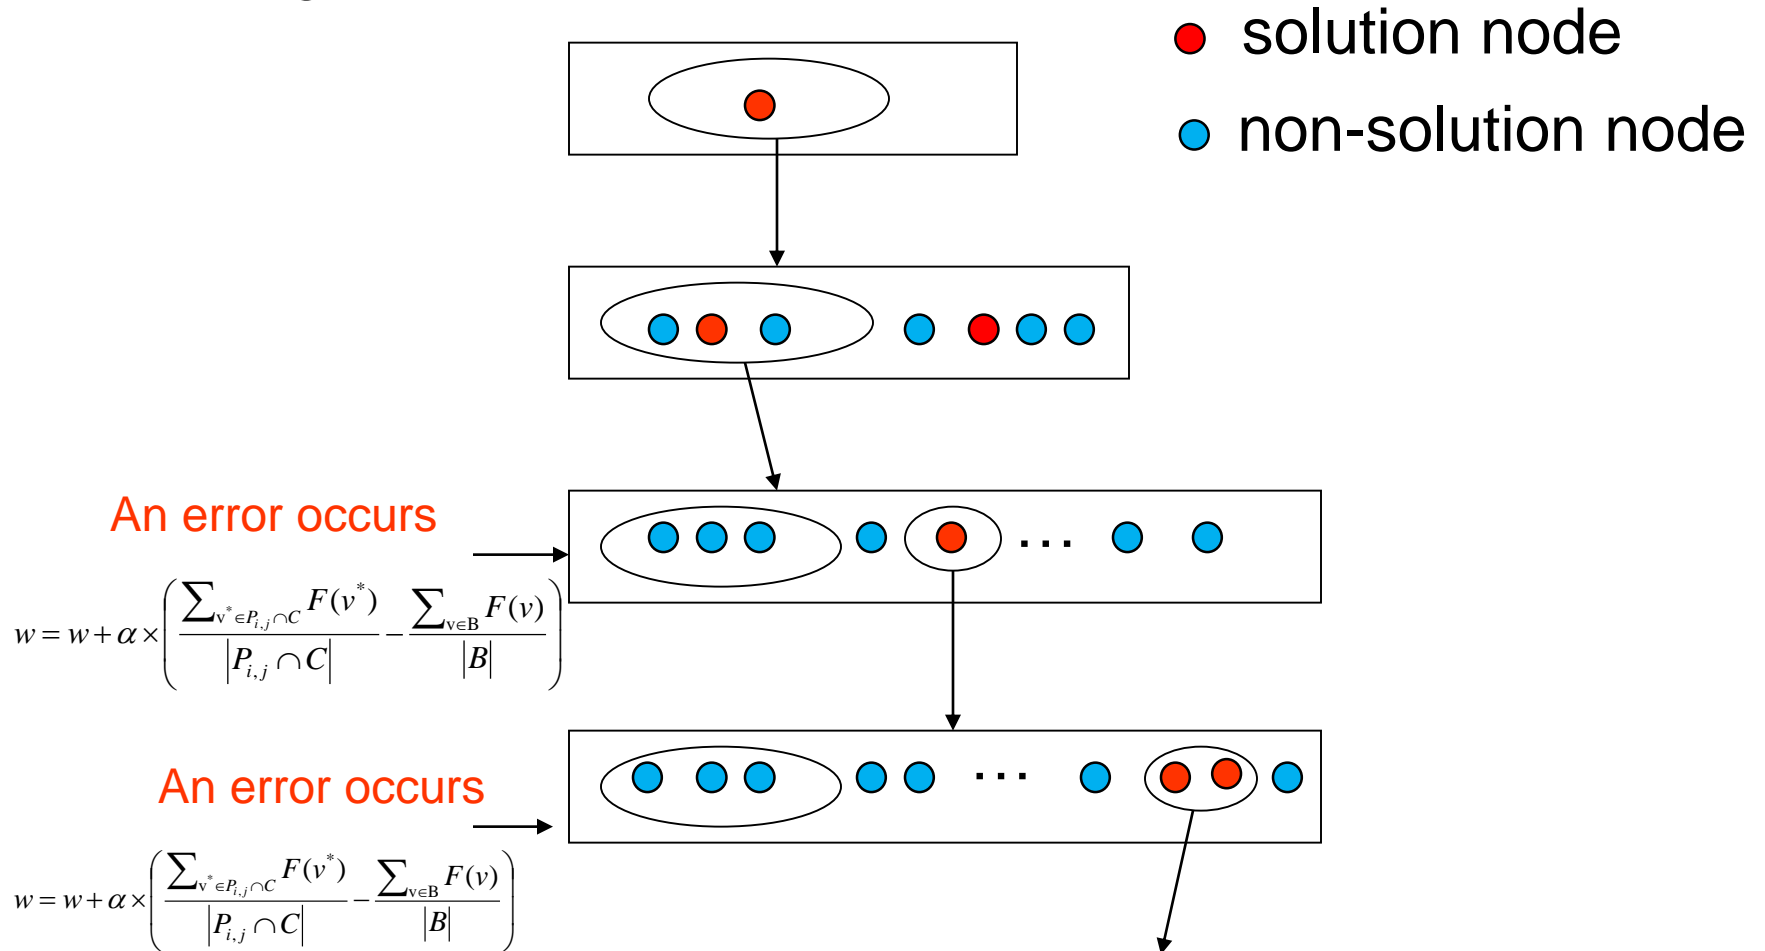


Beam Search Framework: LaSO

- **Search error:** NO target node in the beam
 - ▲ We cannot reach the goal node (correct structured output)
- **Weight update:** perceptron update
 - ▲ $w_{new} = w_{old} + \alpha \cdot (\psi_{avg}(target) - \psi_{avg}(non - target))$
 - ▲ $\psi_{avg}(target)$ = Average features of all target nodes in the candidate set
 - ▲ $\psi_{avg}(non - target)$ = Average features of all non-target nodes in the candidate set
 - ▲ **Intuition:** increase the score of target nodes and decrease the score of the non-target nodes
- **Beam update:** reset beam with target nodes in the candidate set

LaSO Training: Illustration

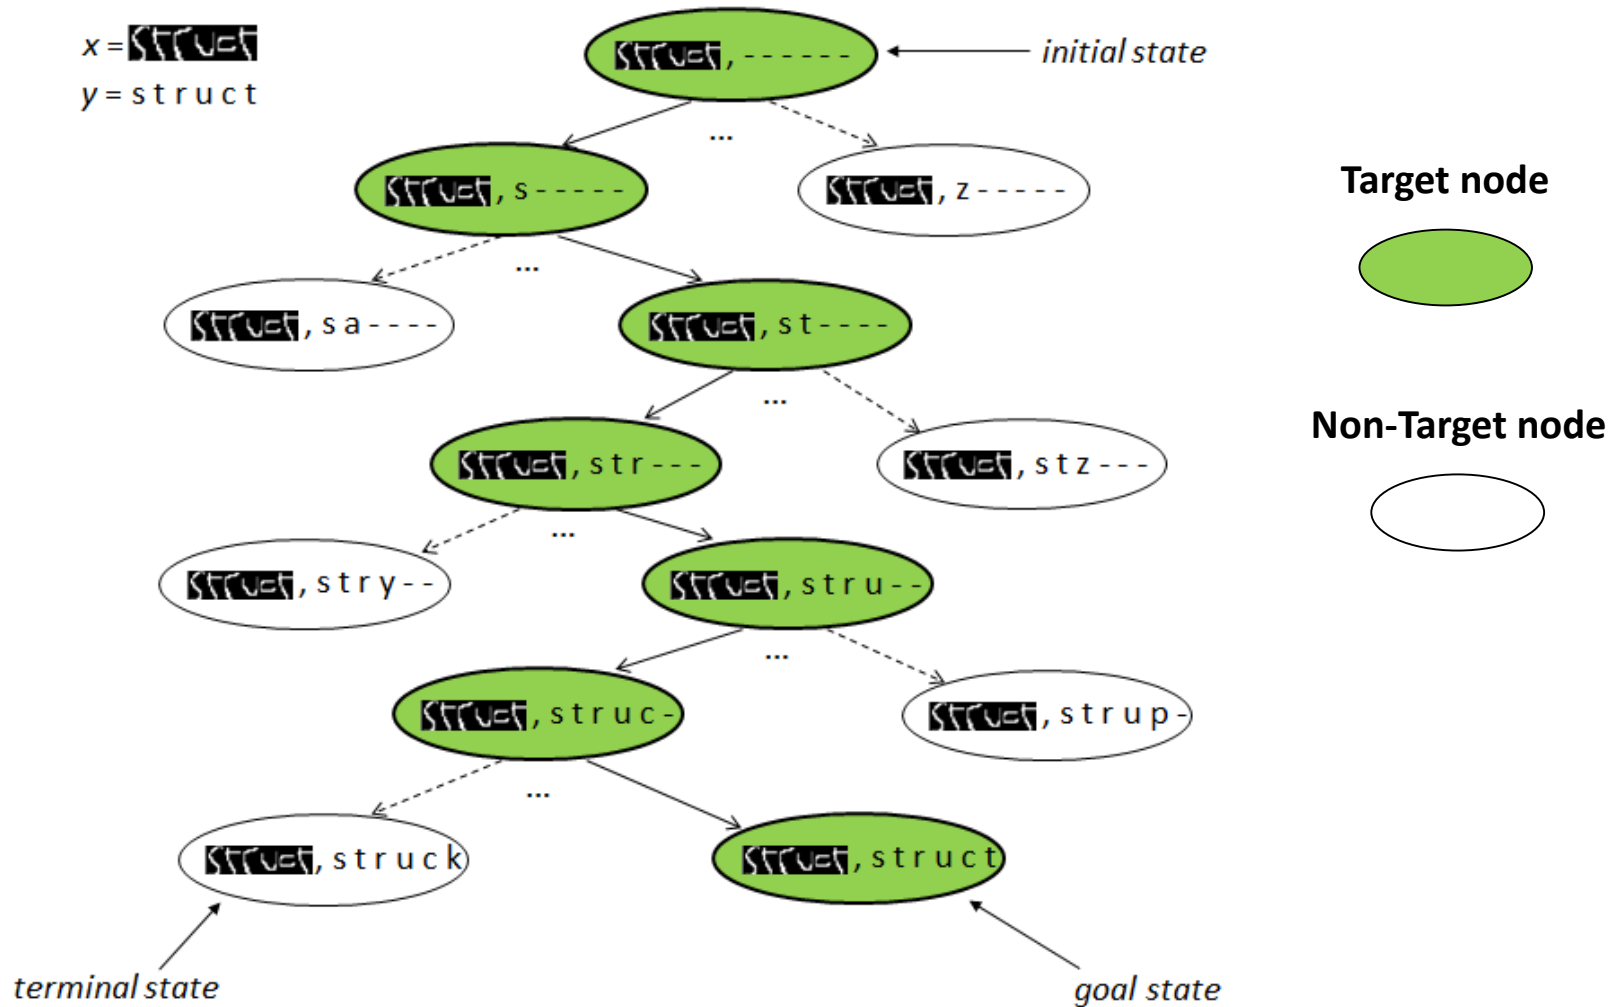
Basic Idea: repeatedly conduct search on training examples
update weights when error occurs



LaSO Framework: Learning

- Represent heuristic function as a linear function
 - ▲ $H(n) = w \cdot \psi(n)$, where $\psi(n)$ stands for features of node n .
- **repeat**
 - ▲ For every training example (x, y)
 - ▲ Perform search with current weights
 - ▲ If **search error** , **update weights**
- **until** *convergence* or *max. iterations*

Search Error: Illustration



Search Error: Illustration for Beam Search

- **Search Error**

- ▶ If the BEAM (internal memory of the search procedure – beam search) **does NOT contain at least one target node**
- ▶ Why?
- ▶ Because we cannot reach the goal node (correct complete output)
- ▶ I call this definition **Conservative**

Search Error: Illustration for Beam Search

- **Search Error**

- ▲ If the BEAM (internal memory of the search procedure – beam search) **contains at least one non-target node ranked higher than a target node**
- ▲ I call this definition **Aggressive**
- ▲ Does not exist in the literature!

- **Conservative vs. Aggressive**

- ▲ ??

Search Error: Illustration for Beam Search

- **Aggressive Error**

- ▶ Seems over-constrained, which may make the learning problem harder than necessary?
- ▶ More training data?
- ▶ Increase the chances of reaching the terminal quickly?

- **Conservative Error**

- ▶ relatively under-constrained. Good or Bad?
- ▶ Less training data?
- ▶ It can keep one target node in the beam and keep on expanding the non-target nodes – no progress?

Search Error: Illustration for Beam Search

- Checking for Search Error (w/ Beam Search)
 - ▲ Compute the candidate set
 - ▲ Compute the score of each node: $H(n) = w \cdot \psi(n)$
 - ▲ Sort the nodes based on the heuristic scores
 - ▲ Select the top scoring B nodes (BEAM)
 - ▲ Apply the search error definition on BEAM
- If **search error** , **update weights**

Weight Update: Illustration for Beam Search

- If **search error**, **update weights**

- **LaSO Weight Update:**

$$\triangle w_{new} = w_{old} + \alpha \cdot (\psi_{avg}(target) - \psi_{avg}(non - target))$$

$\psi_{avg}(target)$ = Average features of all target nodes in the candidate set

$\psi_{avg}(non - target)$ = Average features of all non-target nodes in the candidate set

- **Intuition:** try to increase the score of target nodes and decrease the score of the non-target nodes

What happens after weight update?

- **Re-Score** the target nodes in the candidate set using the **new weights**
- **Reset the BEAM** with the **top-scoring B target nodes**
- Continue the search

Putting Everything Together: LaSO Training

- Initialize weights $w = 0$
- repeat
 - ▲ For every training example (x, y)
 - ▲ Perform search with current weights
 - ▲ If search error, update weights
 - ▲ Re-score the target nodes and reset BEAM
- until *convergence or max. iterations*

Putting Everything Together: LaSO Inference

- **Input:** learned weights w ; structured input x
- **repeat**
 - ▲ Perform search with current weights
- **until** *reaching a terminal state*
- **Output:** the complete output y corresponding to the terminal state

Consistency, Convergence, and Hardness

- **Consistency Problem:**

- ▲ Does a weight vector exist that can solve all the training problems correctly for a given search procedure?

- **Convergence:**

- ▲ Under what definition(s) of margin will LaSO converge?

- **Hardness:**

- ▲ Relating hardness of learning to the size of the beam
- ▲ Intuitively, learning problem becomes easier as we increase the beam size
- ▲ **Degenerate case:** BEAM = infinity, no learning is needed!

Additional Reading

- Yuehua Xu, Alan Fern, and Sungwook Yoon. (2009). **Learning Linear Ranking Functions for Beam Search with Application to Planning.** *Journal of Machine Learning Research (JMLR)*, 10, 1349-1388.
<http://www.jmlr.org/papers/volume10/xu09c/xu09c.pdf>
- Beautiful paper! One of my favorite papers.