

Search

School of EECS
Washington State University

Overview

- ▶ Problem-solving agent
- ▶ Formulating problems
- ▶ Search
- ▶ Uninformed search
- ▶ Informed (heuristic) search
- ▶ Heuristics
- ▶ Admissibility

Problem-Solving Agent

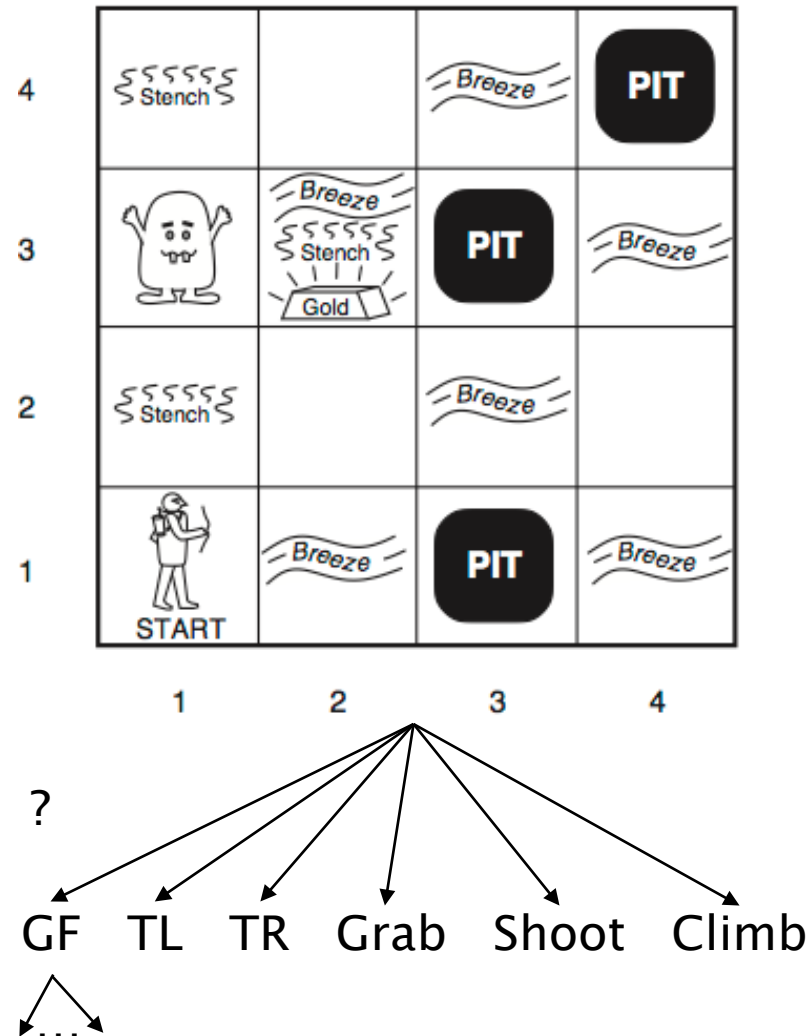
- ▶ Goal-based
- ▶ Atomic state representation
- ▶ Assume solution is a fixed sequence of actions
- ▶ Rationality: Achieve goal (minimize cost)
- ▶ Search for sequence of actions achieving goal

Environment Assumptions

- ▶ **Observable**
 - Agent always knows what state it is in
- ▶ **Deterministic**
 - Each action has one possible outcome (next state)
- ▶ **Discrete**
 - Each state has finite number of applicable actions
- ▶ **Known**
 - Agent knows which state each action will lead to
- ▶ **Once solution sequence known, execute blindly (ignore percepts) until completion**

Wumpus World Example

- ▶ Initial state →
- ▶ Goal state
 - Any state where agent has gold and not in cave
- ▶ Solution?



Problem-Solving Agent

function **SIMPLE-PROBLEM-SOLVING-AGENT** (*percept*) **returns** an action

persistent: *seq*, an action sequence initially empty

state, some description of the current world state

goal, a goal, initially null

problem, a problem formulation

state \leftarrow UPDATE-STATE (*state*, *percept*)

if *seq* is empty **then**

goal = FORMULATE-GOAL (*state*)

problem = FORMULATE-PROBLEM (*state*, *goal*)

seq = SEARCH (*problem*)

if *seq* = *failure* **then return** a null action

action = FIRST (*seq*)

seq = REST (*seq*)

return *action*

Well-Defined Problems

- ▶ State representation (atomic, but...)
- ▶ Action definitions (action: state \rightarrow state)
- ▶ Five parts
 - Initial state
 - Actions
 - Transition model
 - Goal test
 - Path cost

Well-Defined Problems (5 parts)

1. Initial state


2. Actions

- $ACTIONS(s)$ returns set of actions applicable to state s

3. Transition model

- $RESULT(s, a)$ returns state resulting from taking action a in state s
- Successor state is any state reachable from the current state by a single action
- ▶ State space is set of all states reachable from the initial state by any sequence of actions
- ▶ State space forms a directed graph of nodes (states) and edges (actions)
- ▶ Path in state space is a sequence of states connected by actions

Vacuum World State Space



Well-Defined Problems (5 parts)

4. Goal test

- True for any state satisfying goal

5. Path cost

- Sum of the costs of the individual actions along the path
- Step cost $c(s, a, s')$ is the cost of taking action a in state s to reach state s'
- Non-negative
- ▶ Solution is sequence of actions leading from the initial state to a goal state
- ▶ Optimal solution is a solution with minimal path cost



Vacuum World Problem

- ▶ State representation
 - Location of vacuum: Left, Right
 - Cleanliness of each room: Clean, Dirty
 - Example state: (Left,Clean,Clean)
 - How many unique states?
- ▶ **Initial state**: Any state
- ▶ **Actions**: Left, Right, Suck
- ▶ **Transition model**
 - E.g., $\text{Result}((\text{Left}, \text{Dirty}, \text{Clean}), \text{Suck}) = (\text{Left}, \text{Clean}, \text{Clean})$
- ▶ **Goal test**: State = (?,Clean,Clean)
- ▶ **Path cost**
 - Number of actions in solution (step cost = 1)

8-Puzzle

- ▶ State: Location of each tile (and blank)
 - E.g., (B,1,2,3,4,5,6,7,8)
 - How many states?
- ▶ **Initial state**: Any state
- ▶ **Actions**: Move blank Up, Down, Left or Right
- ▶ **Transition model**
- ▶ **Goal test**: State matches Goal State
- ▶ **Path cost**: Number of steps in path (step cost = 1)

7	2	4
5		6
8	3	1

Start State

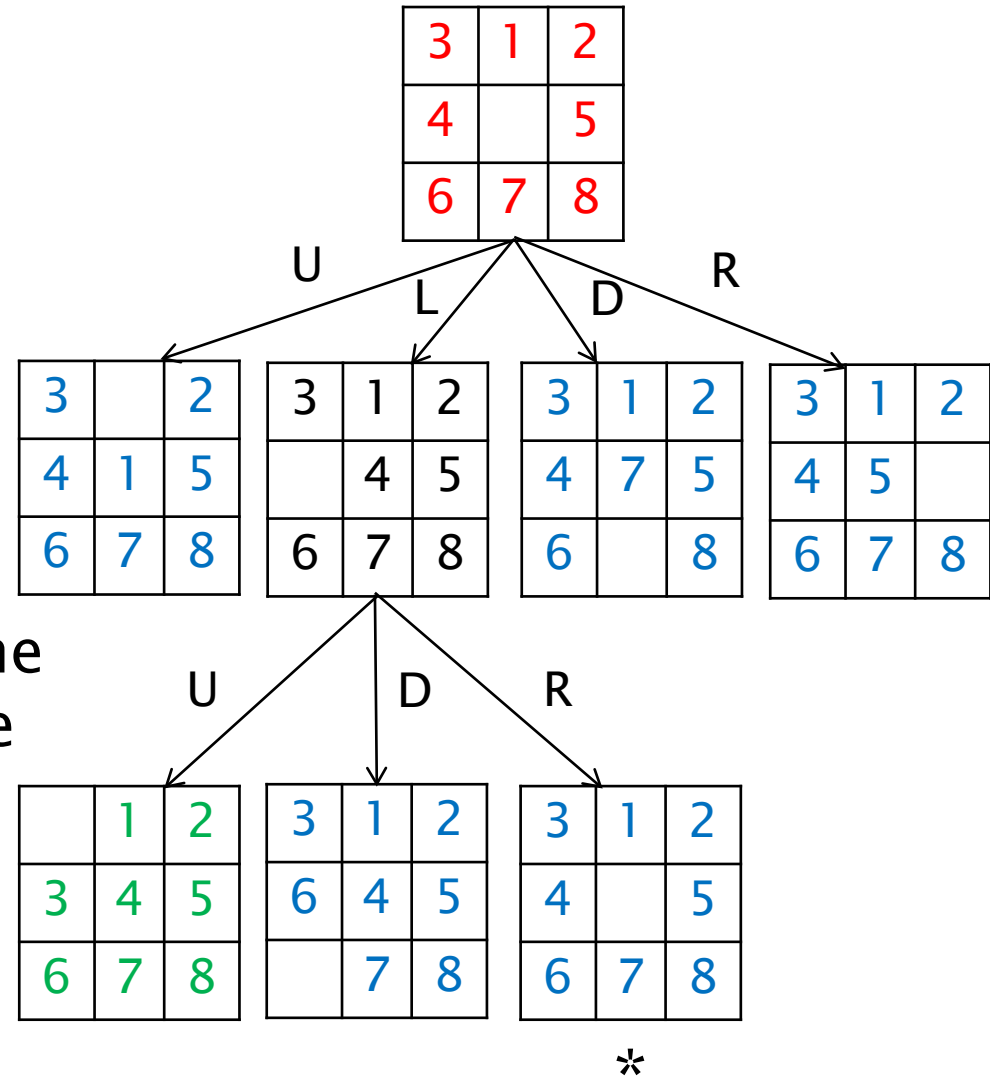
	1	2
3	4	5
6	7	8

Goal State

Search

► Search tree

- Root node is **initial state**
- Node branches for each applicable move from node's state
- **Frontier** consists of the leaf nodes that can be expanded
- Repeated states (*)
- **Goal state**





Search Demo

- ▶ Nice 8-puzzle search web app
 - <http://github.com/tristanpenman/n-puzzle>

1	2	3
	4	5
7	8	6

Initial State

1	2	3
4	5	6
7	8	

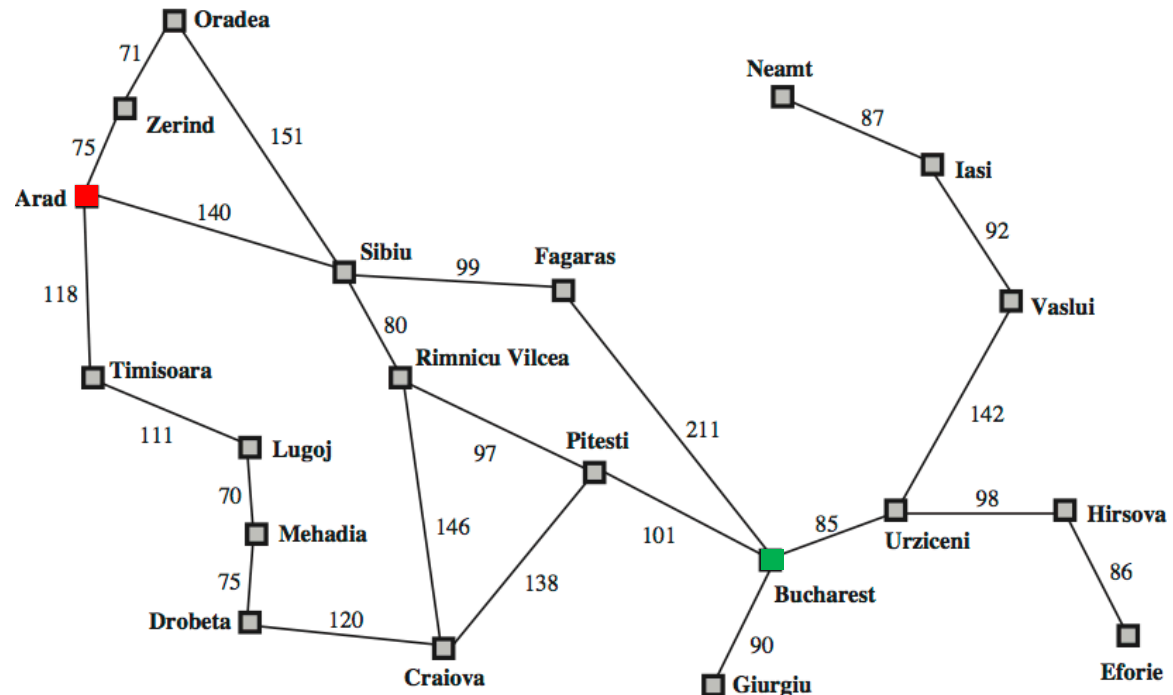
Goal State

Real-World Search Problems

- ▶ Route finding
 - ▶ Touring
 - ▶ Layouts
 - ▶ Robot navigation
 - ▶ Assembly sequencing
 - ▶ Chemical design
-
- ▶ Most of AI can be cast as a search problem

Route Finding Example


- ▶ Romania road map
- ▶ Initial state: **Arad**
- ▶ Goal state: **Bucharest**




Route Finding Example

Search Tree

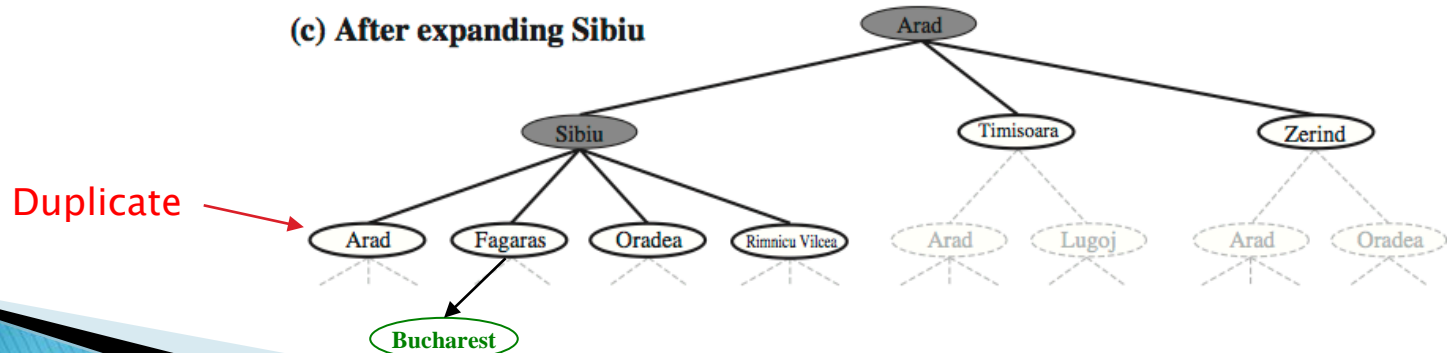
(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu



Tree Search

function **TREE-SEARCH** (*problem*) **returns** a solution, or failure

 initialize the frontier using the initial state of *problem*

loop do

if the frontier is empty **then return** failure

 choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

 expand the node, adding the resulting nodes to the frontier


- ▶ Search strategy determines how nodes are chosen for expansion
- ▶ Suffers from repeated state generation

Graph Search


```
function GRAPH-SEARCH (problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the node, adding the resulting nodes to the frontier
      only if not in the frontier or explored set
```

- ▶ Keep track of explored set to avoid repeated states
- ▶ Changes from TREE-SEARCH highlighted

Graph Search Example



Implementation



```
function CHILD-NODE (problem, parent, action) returns a node
return a node with
    STATE = problem.RESULT (parent.STATE, action),
    PARENT = parent,
    ACTION = action,
    PATH-COST = parent.PATH-COST +
                problem.STEP-COST (parent.STATE, action)
```



Implementation

- ▶ Frontier is a queue or stack
 - How nodes are added/removed defines search strategy
- ▶ Explored set is a hash table
 - Can be large (# unique states)
 - Key is some canonical state representation

Measuring Performance

- ▶ Completeness
 - Is the search algorithm guaranteed to find a solution if one exists?
- ▶ Optimality
 - Does the search algorithm find the optimal solution?
- ▶ Time and space complexity
 - **Branching factor b** (maximum successors of a node)
 - **Depth d** of shallowest goal node
 - **Maximum path length m**
 - Complexity $O(b^d)$ to $O(b^m)$

Uninformed Search Strategies

- ▶ No preference over states based on “closeness” to goal
- ▶ Strategies
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search


Breadth-First Search

- ▶ Expand shallowest nodes in frontier
- ▶ Frontier is a simple queue
 - Dequeue nodes from front, enqueue nodes to back
 - First-In, First-Out (FIFO)

Breadth-First Search

```
function BREADTH-FIRST-SEARCH (problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier ← FIFO queue with node as only element
  explored ← empty set
  loop do
    if EMPTY(frontier) then return failure
    node ← DEQUEUE(frontier) // choose shallowest node in frontier
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child = CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier ← ENQUEUE(child, frontier)
```

Breadth-First Search



► 8-puzzle demo

1	2	3
4	8	5
7		6

Initial State

1	2	3
4	5	6
7	8	

Goal State



Breadth-First Search

- ▶ Complete?
- ▶ Optimal?
- ▶ Time complexity
 - Number of nodes generated (worst case)

$$\sum_{i=0}^d b^i = O(b^{d+1})$$

- ▶ Space complexity
 - $O(b^{d-1})$ nodes in explored set
 - $O(b^d)$ nodes in frontier
 - Total $O(b^d)$

Breadth-First Search

- ▶ Exponential complexity $O(b^d)$
- ▶ For $b=4$, 1KB/node, 1M nodes/sec

Depth	Nodes	Time	Memory
2	16	0.02 ms	16 KB (10^3)
4	256	0.26 ms	256 KB (10^3)
8	65,536	0.07 sec	65 MB (10^6)
16	4.3B	71.6 min	4.3 TB (10^{12})
20	10^{12}	12.7 days	1 PetaByte (10^{15})
30	10^{18}	366 centuries	1 ZettaByte (10^{21})

Uniform-Cost Search

- ▶ Expand node n with lowest path cost $g(n)$
- ▶ Frontier is a priority queue
 - Queue partially ordered by path cost
 - Lowest path cost node always at the front

Uniform-Cost Search


```
function UNIFORM-COST-SEARCH (problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier ← priority queue ordered by PATH-COST, with node as only element
  explored ← empty set
  loop do
    if EMPTY(frontier) then return failure
    node ← DEQUEUE(frontier) // choose lowest cost node in frontier
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child = CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier ← ENQUEUE(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child
```

Why not
check
Goal-Test
here?

Why is this test
necessary?

Uniform-Cost Search

- ▶ Example (Sibiu → Bucharest)





Uniform-Cost Search


- ▶ Complete?
- ▶ Optimal?
- ▶ Time and space complexity
 - b = branching factor
 - ε = minimum step cost (>0)
 - C^* = cost of optimal solution

$$O(b^{1+\lfloor C^*/\varepsilon \rfloor})$$

Depth–First Search

- ▶ Always expand the deepest node
- ▶ Frontier is a simple queue
 - Enqueue nodes to front, dequeue nodes from front
 - Last–In, First–Out (LIFO)
- ▶ Otherwise, same code as BFS
- ▶ Or, implement recursively

Depth-First Search



DEMO



Depth-First Search

- ▶ Tree-Search version
 - Not complete (infinite loops)
 - Not optimal
- ▶ Graph-Search version
 - Complete
 - Not optimal
- ▶ Time complexity (m = max depth): $O(b^m)$
- ▶ Space complexity
 - Tree-search: $O(bm)$
 - Graph-search: $O(b^m)$

Depth-Limited Search

function **DEPTH-LIMITED-SEARCH** (*problem, limit*) **returns** a solution, or failure/cutoff
return RECURSIVE-DLS (MAKE-NODE (*problem*.INITIAL-STATE), *problem, limit*)

function **RECURSIVE-DLS** (*node, problem, limit*) **returns** a solution, or failure/cutoff
if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)
else if *limit* = 0 **then return** *cutoff*
else
 cutoff_occurred \leftarrow false
 for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**
 child = CHILD-NODE(*problem, node, action*)
 result \leftarrow RECURSIVE-DLS (*child, problem, limit - 1*)
 if *result* = *cutoff* **then** *cutoff_occurred* \leftarrow true
 else if *result* \neq failure **then return** *result*
 if *cutoff_occurred* **then return** *cutoff* **else return** failure

Depth-Limited Search

- ▶ Limit DFS depth to ℓ
- ▶ Still incomplete, if $\ell < d$
- ▶ Non-optimal if $\ell > d$
- ▶ Time complexity: $O(b^\ell)$
- ▶ Space complexity: $O(b\ell)$

Iterative-Deepening Search

- ▶ Run DEPTH-LIMITED-SEARCH iteratively with increasing depth limit

```
function ITERATIVE-DEEPENING-SEARCH (problem) returns a solution, or failure  
  for depth = 0 to  $\infty$  do  
    result = DEPTH-LIMITED-SEARCH (problem, depth)  
    if result  $\neq$  cutoff then return result
```

Iterative-Deepening Search


Limit = 0




Limit = 1



Limit = 2



Limit = 3



DEMO



Iterative-Deepening Search


- ▶ Complete?
- ▶ Optimal?
- ▶ Space complexity: $O(bd)$
- ▶ Time complexity

$$\sum_{i=0}^{d-1} (d-i)b^{i+1} = (d)b + (d-1)b^2 + \dots + (1)b^d = O(b^d)$$

- Nodes at depth d = all nodes at depths 1 to $(d-1)$
- ▶ Iterative deepening best uninformed search when solution depth unknown

Bidirectional Search

- ▶ Search forward from initial state and backward from goal state
- ▶ Meet (hopefully) in the middle
- ▶ Each search has complexity $O(b^{d/2}) \ll O(b^d)$
- ▶ Replace goal test with frontier intersection
- ▶ How to reverse actions?



Uninformed Search Strategies

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional
Complete	Yes ¹	Yes ^{1,2}	No	No	Yes ¹	Yes ^{1,4}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\varepsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\varepsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal	Yes ³	Yes	No	No	Yes ³	Yes ^{3,4}

1. Complete if b is finite
2. Complete if step costs $\geq \varepsilon > 0$
3. Optimal if step costs all the same
4. If both directions use BFS

Informed (Heuristic) Search

- ▶ Guided by problem-specific knowledge other than the problem formulation
- ▶ Problem-specific knowledge usually expressed as heuristics

Heuristic Function

- ▶ Heuristic function $h(n)$ estimates cost of the path from state n to a goal state
 - E.g., 8-puzzle
 - Number of tiles out
 - Euclidean distance of each tile
 - City-block (Manhattan) distance of each tile
 - Non-negative function
 - For goal node $h(n)=0$
- ▶ Recall path cost $g(n)$ is the cost so far from the initial state to state n
- ▶ Evaluation function $f(n) = g(n) + h(n)$ estimates the total cost of a solution going through state n

1	5	2
4	3	
7	8	6

1	2	3
4	5	6
7	8	

Goal State



Best-First Search

- ▶ Choose next frontier node with smallest $f(n)$
- ▶ Depth-first search = Best-first search with
 - $f(n) = g(n) + h(n) = ?$
- ▶ Breadth-first search = Best-first search with
 - $f(n) = g(n) + h(n) = ?$
- ▶ Uniform-cost search = Best-first search with
 - $f(n) = g(n) + h(n) = ?$

Heuristic Search Strategies

- ▶ Greedy best-first search
- ▶ A* search
- ▶ Memory-bounded heuristic search

Greedy Best-First Search

- ▶ Best-first search with $f(n) = h(n)$
- ▶ Example: Route-finding problem
 - $h(n)$ = straight-line distance from city n to goal city

Straight-line distances to Bucharest:


Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy Best-First Search Example: Arad to Bucharest

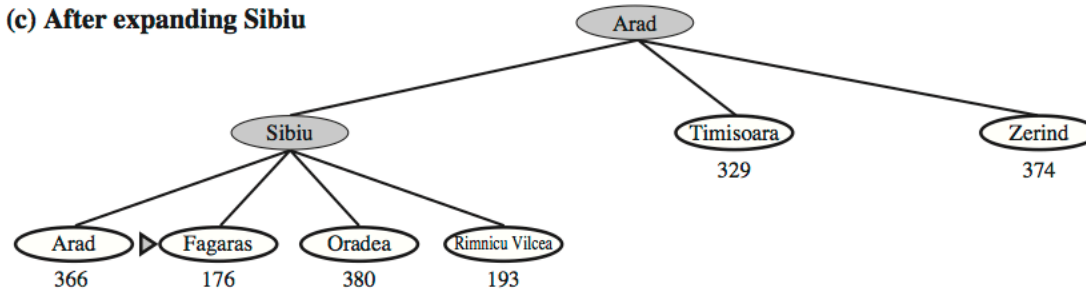
(a) The initial state



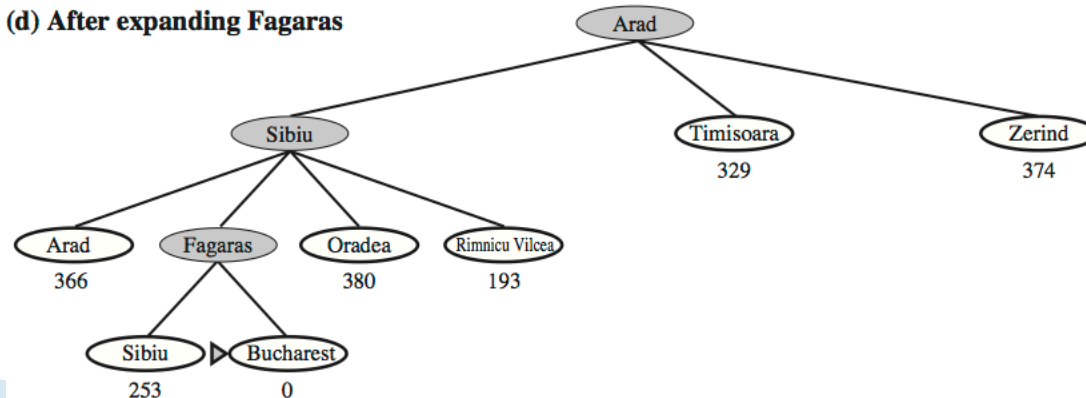
(b) After expanding Arad



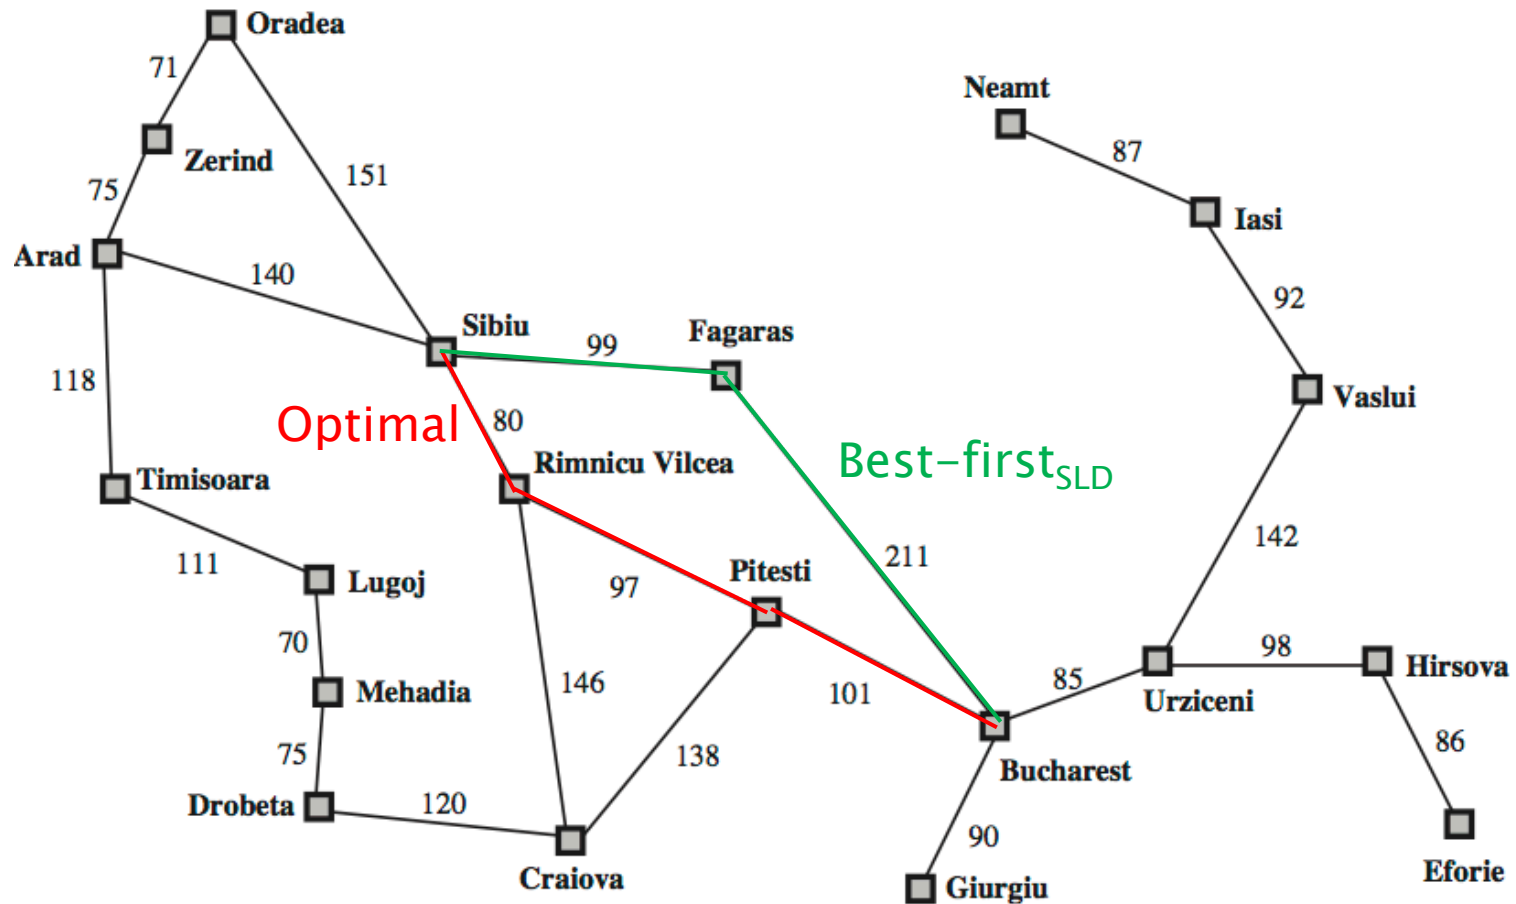
(c) After expanding Sibiu



(d) After expanding Fagaras



Greedy Best-First Search Example: Arad to Bucharest



Greedy Best-First **Tree** Search

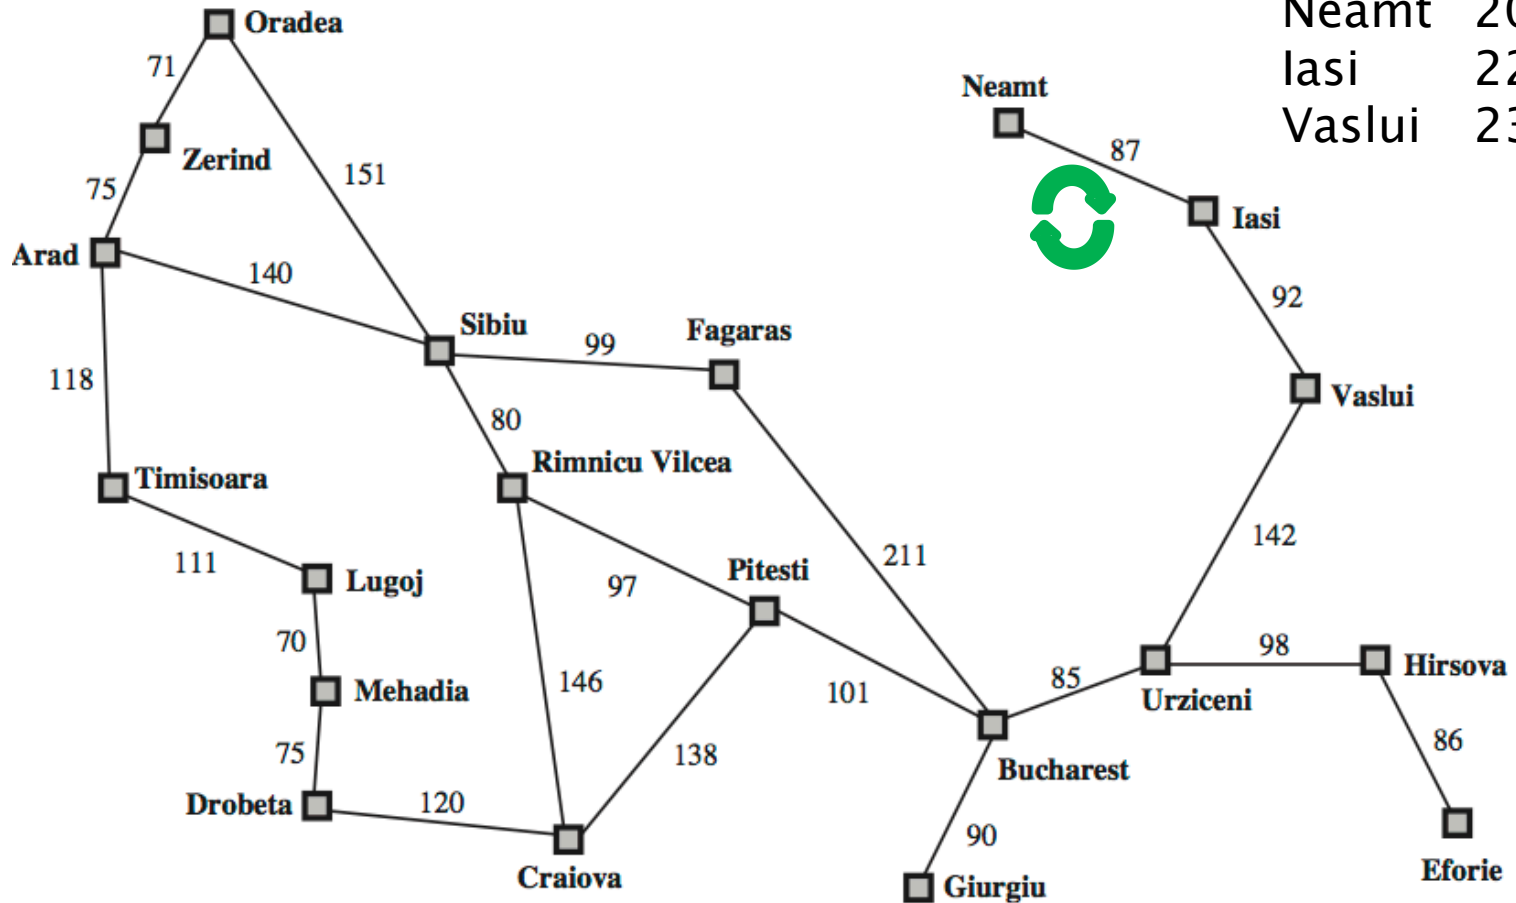
Example: Iasi to Fagaras

SLD to Fagaras

Neamt 200

Iasi 220

Vaslui 230





Greedy Best-First Search

- ▶ Complete?
- ▶ Optimal?
- ▶ Time and space complexity: $O(b^m)$
 - b = branching factor
 - m = maximum depth of search space
 - Worst case
 - Good heuristic can substantially improve

DEMO


A* Search

History: A* generalizes over algorithms A1 and A2, which were heuristic extensions to Dijkstra's shortest path algorithm.


- ▶ $f(n) = g(n) + h(n)$
 - Estimated cost of solution through n
- ▶ Same as Uniform-Cost search using $f(n)$
- ▶ Complete and optimal under some constraints on $h(n)$
- ▶ Example: Route-finding using SLD

A* Search Example: Arad to Bucharest


(a) The initial state



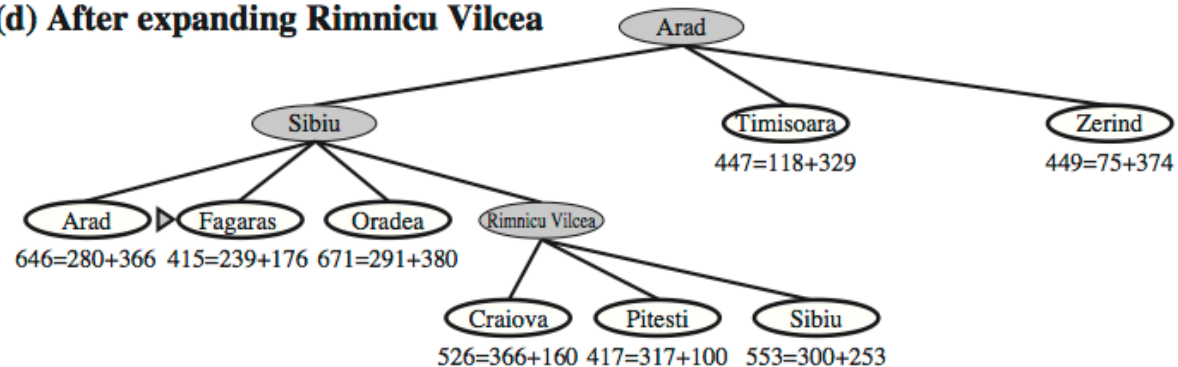
(b) After expanding Arad



(c) After expanding Sibiu

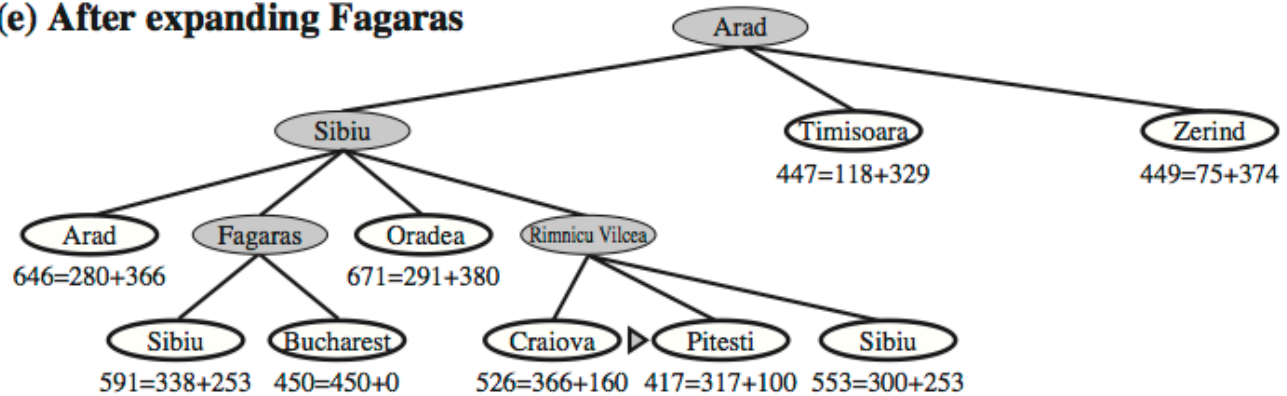


(d) After expanding Rimnicu Vilcea

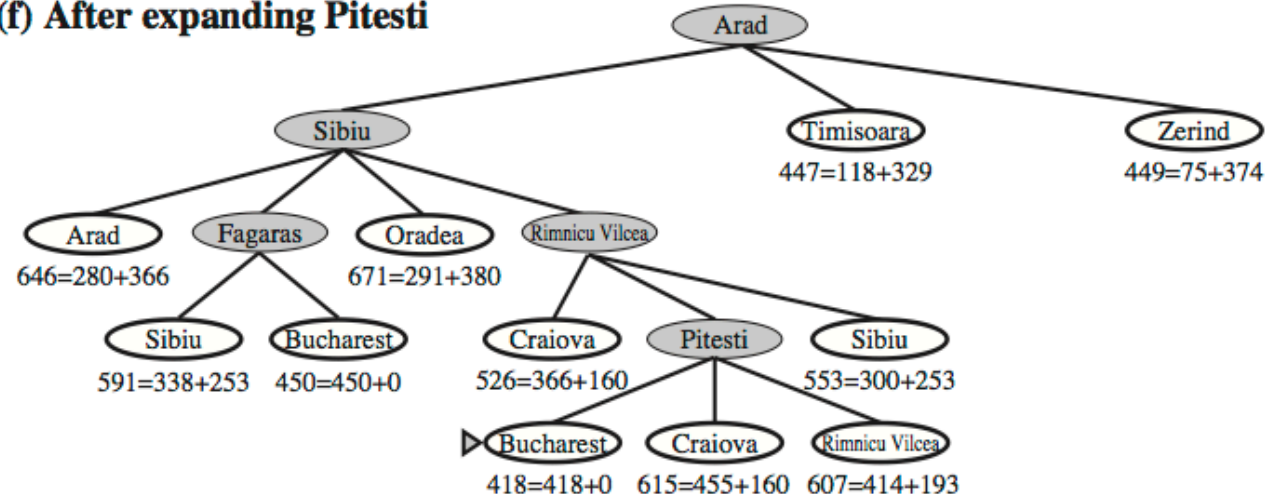


A* Search Example: Arad to Bucharest (cont.)

(e) After expanding Fagaras



(f) After expanding Pitesti





Optimality of A*

- ▶ For A* **tree** search to be optimal, $h(n)$ must be admissible
 - A heuristic function $h(n)$ is admissible if it never over-estimates the cost of reaching the goal from n
 - E.g., Straight-line distance for route finding
 - E.g., Tiles out of place in 8-puzzle
- ▶ For A* **graph** search to be optimal, heuristic must further satisfy triangle inequality (also called consistent or monotonic)
 - A heuristic function $h(n)$ satisfies the triangle inequality if $h(n) \leq \text{cost}(n, a, n') + h(n')$

A* Search

- ▶ Complete and optimal?
 - Yes, if heuristic is admissible
- ▶ Time and space complexity?
 - Still $O(b^d)$ worst case
 - Space is typically the bottleneck
- ▶ A* is optimally efficient
 - No other algorithm using the same consistent heuristic is guaranteed to expand fewer nodes

DEMO

Memory-Bounded Heuristic Search

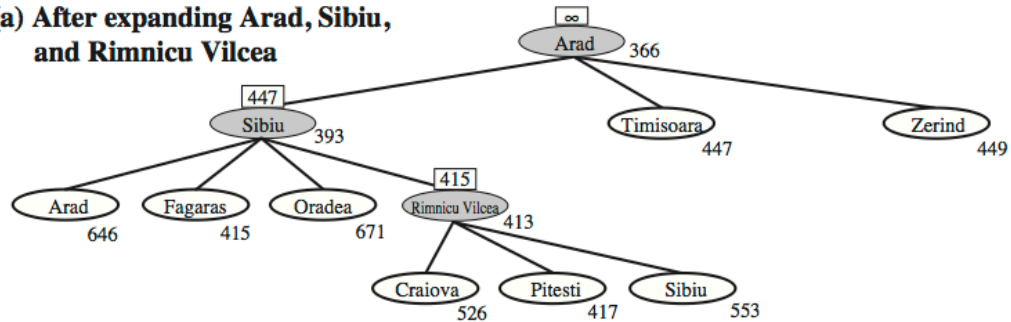
- ▶ Iterative-Deepening A^* (IDA*)
 - Like iterative deepening, except uses limit on f rather than depth
 - Next f limit is the smallest f of a node exceeding the limit in the previous iteration
 - Space efficient, but may take long if f values increase slowly

Memory-Bounded Heuristic Search

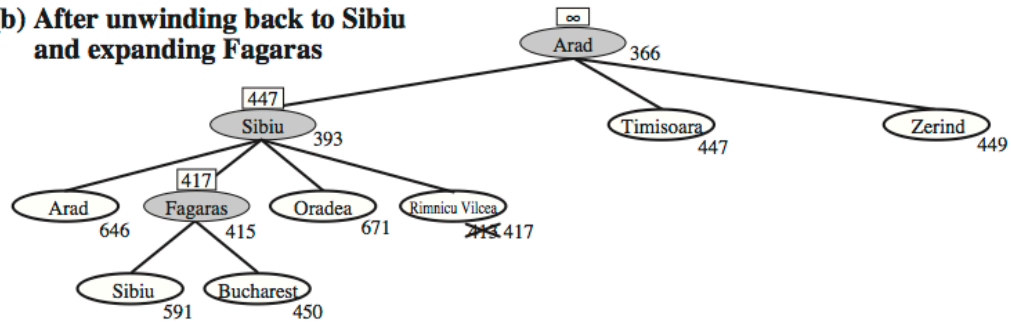
- ▶ Recursive Best-First Search (RBFS)
 - Similar to recursive depth-first search
 - Each node along current path maintains best f value (f -limit) of alternative path from an ancestor
 - If current node's f value exceeds f -limit then backtrack to ancestor and expand alternative path

RBFS Example

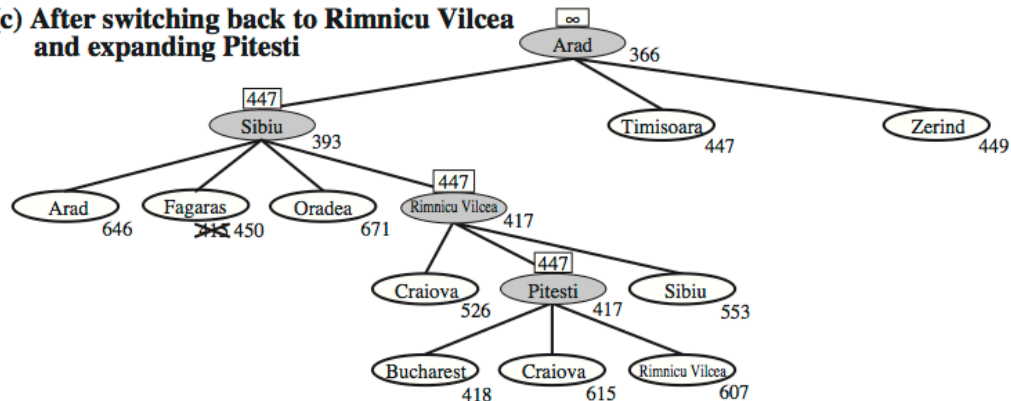
(a) After expanding Arad, Sibiu, and Rimnicu Vilcea



(b) After unwinding back to Sibiu and expanding Fagaras



(c) After switching back to Rimnicu Vilcea and expanding Pitesti



Memory-Bounded Heuristic Search

- ▶ Recursive Best-First Search (RBFS)
 - Like IDA*, RBFS explores same states many times
 - IDA* and RBFS use only linear memory
 - They cannot take advantage of more memory, if available

Memory-Bounded Heuristic Search

- ▶ Simplified Memory-bounded A^* (SMA*)
 - Similar to A^*
 - Nodes maintain best f -value of any explored node in their subtree
 - If out of memory, remove node with worst f -value and update parent's f -value
 - Complete and optimal if best solution is reachable within memory



Heuristics Revisited

- ▶ Why not use $h(n) = 1$?
- ▶ How to measure quality of heuristic?
- ▶ Effective branching factor b^*
 - Assume A^* generates N nodes to find solution at depth d
 - What branching factor needed for a uniform tree of depth d to include $N+1$ nodes?
 - $N + 1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$
 - Ideally, $b^* = 1$
- ▶ E.g., $N=52$, $d=5$, $b^*=1.92$

Heuristics Revisited

- ▶ E.g., 8-puzzle
 - h_1 = tiles out of place
 - h_2 = sum of tiles' city block distances

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

$$h_1 = 8$$

$$h_2 = 3 + 1 + 2 + 2 + 3 + 2 + 2 + 3 = 18$$

$$\text{Solution cost} = 26$$

Heuristics Revisited

- ▶ Values averaged over 100 8-puzzle problems for each d

- ▶ Note:
 $b^*(h_2) \leq b^*(h_1)$

d	Search Cost (nodes generated)			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Heuristics Revisited

- ▶ Heuristic h_2 dominates h_1 if, for all nodes n , $h_2(n) \geq h_1(n)$
- ▶ Implies A^* using h_2 will typically generate fewer nodes than A^* using h_1
- ▶ “City block distance” dominates “misplaced tiles”
- ▶ In general, want $h(n)$ to be consistent and close to true solution cost from node n
 - But still be fast to compute

Designing Heuristics

- ▶ Relaxed problems
 - $h(n)$ = cost of solution to relaxed problem
 - E.g., 8-puzzle where you can swap tiles
- ▶ Subproblems
 - $h(n)$ = cost of solution to subproblem
 - E.g., get half the tiles in correct position
- ▶ Learning from experience
 - Collect experience as (state, solution cost) pairs
 - Learn $h(n)$: state \rightarrow solution cost

Summary

- ▶ Problem-solving agent
- ▶ Formulating problems
- ▶ Search
- ▶ Uninformed search (Iterative-Deepening)
- ▶ Informed (heuristic) search (A^*)
- ▶ Admissible heuristics