

Lecture #9: Ensemble Learning

What is Ensemble Learning?

In ensemble learning, the idea is to combine multiple classifiers into a single one. Ensemble learning usually works very well in practice.

Two methods (for this class):

1. Bagging
2. Boosting

Bagging

BAGGING: (Bootstrap AGGREGATING)

1. Input: n labelled training examples $(x_1, y_1), \dots, (x_n, y_n)$

2. Algorithm:

Repeat k times:

(a) Select m samples out of n with replacement from the training set to get training set S_i

(b) Train classifier h_i on S_i (usually, h_i 's are the same type of classifier)

3. Output: Classifiers h_1, \dots, h_k

Testing: Given test example x , output the majority of $h_1(x), h_2(x), \dots, h_k(x)$ (break ties at random as usual)

Choice Points in Bagging

1. How to pick k ?

Higher k is better, but also increases training time, storage requirement and classification time. So pick a k which is feasible.

Choice Points in Bagging

2. How to pick m ?

Popular choice for $m = n$.

But this is still very different from working with the entire training set!

$$\Pr(S_i = S) = \frac{n!}{n^n} \quad \left(\begin{array}{l} \text{\# ways of choosing } n \text{ samples} \\ \text{with replacement} = n^n \end{array} \right)$$

\Rightarrow Only $n!$ of these ways give you the entire training set!

$\Rightarrow \frac{n!}{n^n} \approx$ a very tiny number $\ll 2^{-n/2}$

For any (x_j, y_j) , $\Pr((x_j, y_j) \text{ is not in } S_i) = \left(1 - \frac{1}{n}\right)^n \approx \frac{1}{e}$ (for large n)
 $1/e \approx 0.37$; so about 37% of the data is left out of S_i .

Why does Bagging work?

It can be shown that bagging decreases the variance of a classifier. (it doesn't help much with bias).

Thus it prevents overfitting.

Boosting

Boosting

Sometimes it is:

- easy to come up with simple, easy to use, rules of thumb classifiers
- but hard to come up with a single highly accurate rule.

Examples:


(1) Spam classification, based on email text.

Certain words, eg. "Nigeria", "Online Pharmacy", etc. typically are a good indicator of spam.

Rule-of-thumb: Does email contain word "Nigeria" ?

Boosting (contd.)

(2) Detect if an image has a face in it.



On an average, pixels around the eyes are darker than those below.

Rule of thumb: Is the (average darkness in the shaded region) - (average darkness in the white rectangular region below) > 0 ?

Boosting gives us a way to combine these ~~weak~~ rules ~~into~~ of thumb into good classifiers.

Definitions:

1. Weak Learner: A simple rule of thumb that doesn't necessarily work very well.
2. Strong Learner: A good classifier (with high accuracy)

Boosting Framework

Boosting Procedure:

1. Design method to find a good rule of thumb.
2. Repeat:
 - Find a good rule of thumb
 - Modify training data to get a second data set
 - Apply method ~~of~~ to new data set to get a good rule of thumb, and so on.
1. How to get a good rule of thumb? Application specific (more later)
2. How to modify training data set?
 - Give highest weight to the hardest examples - those that were misclassified more often by previous rules of thumb.
3. How to combine the rules of thumb into a prediction rule?
Take a weighted majority of the rules.

Weak Learner: Example


Let \mathcal{D} be a distribution over labelled examples, and let h be a classifier.

Error of h wrt \mathcal{D} is:

$$\text{err}_{\mathcal{D}}(h) = \Pr_{(x,y) \sim \mathcal{D}} [h(x) \neq y]$$

Example: \mathcal{D} :

X : takes values $\frac{1}{4}, \frac{1}{2}, \frac{3}{4}, 1$, each w.p. $\frac{1}{4}$.



$Y = 1$ if X has ~~the~~ value $> 1/2$, o/w $Y = 0$.

Then if h is the rule:

$$\begin{aligned} h(x) &= 1 \text{ if } x > \frac{1}{4} \\ &= 0 \text{ o/w.} \end{aligned}$$

$$\text{Then, } \text{err}_{\mathcal{D}}(h) = \frac{1}{4}.$$

Basic Definitions

→ h is called a weak learner if $\text{err}_D(h) < 0.5$

→ Error of random guessing is 0.5 (with 2 labels)

Given training examples $(x_1, y_1), \dots, (x_n, y_n)$, we can assign weights w_1, \dots, w_n to these examples. If $\sum_{i=1}^n w_i = 1$, $w_i \geq 0$, we can think of these weights as a probability distribution over the examples.

Error of a classifier h wrt W is:

$$\text{err}_W(h) = \sum_{i=1}^n w_i \mathbf{1}(h(x_i) \neq y_i)$$

$\mathbf{1}$ is the indicator function,
where $\mathbf{1}(P) = 1$ if P is true
 $= 0$ otherwise.

Boosting Algorithm

Input: Training set $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$, $y_i = \pm 1$

$$D_1(i) = 1/n \text{ for all } i = 1, \dots, n$$

For $t = 1, 2, 3, \dots$

$h_t = \text{weak-learner wrt } D_t$. (so, $\text{err}_{D_t}(h_t) < 0.5$)

$$\epsilon_t = \text{err}_{D_t}(h_t)$$

$$\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t} \quad \left(\text{so, } \alpha_t \text{ is high when } \epsilon_t \text{ is low,} \right. \\ \left. \text{and almost 0 when } \epsilon_t \text{ is close to 0.5} \right)$$

$$D_{t+1}(i) = \frac{D_t(i) e^{-\alpha_t y_i h_t(x_i)}}{Z_t} \quad \left(D_{t+1} \text{ goes } \uparrow \text{ if } i \text{ is misclassified} \right. \\ \left. \text{by } h_t; \text{ so higher } D_t \text{ means} \right. \\ \left. \text{harder example.} \right)$$

where Z_t is a normalization constant to ensure that

$$\sum_i D_{t+1}(i) = 1.$$

Final classifier: $H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$ (weighted majority
of $h_t(x)$'s)

Boosting Examples

Example of Weighted Error:

Suppose training data is: $((0,0), 1), ((1,0), 1), ((0,1), -1)$

weights W : $\frac{1}{2}$ $\frac{1}{4}$ $\frac{1}{4}$

classification rule: Predict 1 if $x_1 \leq \frac{1}{2}$, -1 otherwise.

$$\text{err}_w(h) = \frac{1}{2} \times 0 + \frac{1}{4} \times 1 + \frac{1}{4} \times 1 = \frac{1}{2}$$


(The usual (unweighted) error would be ~~2/3~~ 2/3).

Boosting Algorithm Example:

Training data:

$((1,1), +)$	$((2,1), -)$	$((4,1), -)$
$((1,2), +)$	$((2,2), -)$	$((3,2), -)$
$((2,3), +)$	$((3,3), +)$	$((4,3), -)$
$((2,4), +)$		

Boosting Examples



Initially: $\mathcal{D}_1(i) = 0.1$ (for all i)

~~Suppose~~ \mathcal{H}

Weak Learners: Set of vertical and horizontal thresholds.

- ① Suppose we pick $h_1(x) = +$ if $x_1 \leq 1.5$
 $= -$ otherwise

Name the points: a, b, \dots, j (for ease of understanding)

Then:

$$\text{err}_{\mathcal{D}_1}(h_1) = \epsilon_1 = 0.3 \quad \alpha_1 = 0.42$$


Weights of a, b, c, d, e, f, g : $\mathcal{D}_2 = 0.07$

Weights of h, i, j : $\mathcal{D}_2 = 0.17$

$$\begin{aligned} Z_2 &= 7 \cdot e^{-0.42 \cdot 0.1} + 3 \cdot 0.1 \cdot e^{0.42} \\ &= 0.92 \end{aligned}$$

Note: Calculations rounded to 2 decimal places.

Boosting Examples



In Round 2, suppose we pick

$$h_2(x) = + \text{ if } x_2 > 2.5 \\ = - \text{ otherwise.}$$

$$\text{err}_{D_2}(h_2) = \epsilon_2 = 0.21$$

$$\alpha_2 = 0.66$$

Weights of a, b: $D_3 := 0.07 \times e^{0.66} / Z_3 = 0.17$

Weights of c, d, e, f: $D_3 := 0.07 \times e^{-0.66} / Z_3 = 0.04$


Weights of h, i, j: $D_3 := 0.17 \times e^{-0.66} / Z_3 = 0.11$

Weight of g: $D_3 := 0.07 \times e^{0.66} / Z_3 = 0.17$

$$Z_3 = 0.81$$

Boosting Examples

③



In Round 3, suppose we pick:

$$h_3(x) = + \text{ if } x_1 \leq 3.5 \\ = - \text{ otherwise.}$$

$$\text{err}_{D_3}(h_3) = \epsilon_3 = 0.12$$

$$\alpha_3 = 0.99$$

Weights of a, b: $D_4 := 0.17 \times e^{-0.99} / Z_4 = 0.1$

" " c, d, e: $D_4 := \cancel{0.17} \times e^{0.99} / Z_4 = 0.04 e^{0.99} / Z_4 = 0.17$

" " h, i, j: $D_4 := 0.11 \times e^{-0.99} / Z_4 = 0.06$

" " f: $D_4 := 0.04 e^{-0.99} / Z_4 = 0.02 \quad Z_4 = 0.65$

" " g: $D_4 := 0.17 e^{-0.99} / Z_4 = 0.1$

Final classifier: $\text{sign}(\alpha_1 h_1(x) + \alpha_2 h_2(x) + \alpha_3 h_3(x))$

$$= \text{sign}(0.42 h_1(x) + 0.66 h_2(x) + 0.99 h_3(x))$$

Boosting and Overfitting


When to stop boosting? Use a validation dataset to find a stopping time.

Stop when validation error does not improve.

Boosting and Overfitting:

Overfitting can happen with boosting, but often does not.


Typical boosting run:



Reason is that the margin of classification often increases with boosting.

Boosting Margin

Intuitively, margin of classification measures how far the + labels are from the - labels.



Note: Notion of margin for boosting is a little different from the exact way we defined margin for perceptron, but the difference is fairly technical.

For boosting:

- think of each $h_t()$ as a feature
- Feature space is:

$$[h_1(x), h_2(x), \dots, h_T(x)]$$

- Margin of example x is: $\left| \sum_{t=1}^T \alpha_t h_t(x) \right|$.
- If you have large margin data, then classifiers need less training examples to avoid overfitting. (This is also why kernels work, even if they are very high dimensional feature spaces.)