

# Lecture #12: Unsupervised Learning

# Supervised vs. Unsupervised Learning

- **Supervised Learning**



- ▲ Training examples are labeled with the class labels

- **Unsupervised Learning**


- ▲ We are only provided with examples without specifying the class labels

# What can we learn from unlabeled data?

- Group of clusters in the data
- Low dimensional structure




PCA



Nonlinear embedding


# Clustering



- Are there any groups in the data?
- How to group?
- How many groups?


# Unsupervised Learning

- **Clustering:** most common form of unsupervised learning
  - ▲ Given a collection of **unlabeled** examples (objects), discover self-similar groups in the data



# Unsupervised Learning

- **Text Clustering**



# Unsupervised Learning

- Image Segmentation



# Clustering Applications

- Find genes that are similar in their functions
- Group documents based on topics
- Categorize customers based on their buying habit
- Group images based on their contents



# Clustering Issues

- What is a natural **grouping** among these objects?
  - ▲ Definition of "group"
- What makes objects "**related**"?
  - ▲ Definition of "similarity/distance"
- **Representation** for objects
  - ▲ Vector, normalization?
- **How many clusters?**
  - ▲ Fixed a priori?
  - ▲ Completely data driven?
  - ▲ Avoid "trivial" clusters - too large or small

# What is a natural grouping?



- By color? By pattern? By weight?
- The definition of natural grouping is subjective
- This is why we call clustering **exploratory** data analysis

# What is similarity?

- This is a philosophical question. We will take a more pragmatic approach.
  - ▲ Depends on representation and algorithm. For many representations/algorithms, it is easier to think in terms of a distance (rather than similarity) between vectors



Hard to define but

We know it when we see it

# Properties of a distance measure?

- **$D$  must be Symmetric**

- ▶  $D(A, B) = D(B, A)$

- ▶ Otherwise, we can say  $A$  looks like  $B$  but  $B$  does not look like  $A$

- **Positivity, and self-similarity**

- ▶  $D(A, B) \geq 0$ , and  $D(A, B) = 0$  iff  $A = B$

- ▶ Otherwise, there will different objects that we cannot tell apart

- **Must satisfy triangle inequality**

- ▶  $D(A, B) + D(B, C) \geq D(A, C)$

- ▶ Otherwise, one can say “ $A$  is like  $B$ ,  $B$  is like  $C$ , but  $A$  is not like  $C$  at all”

# Distance Measures: Minkowski Metric

- Suppose two object  $x$  and  $y$  both have  $d$  features
  - ▲  $x = (x_1, \dots, x_d), y = (y_1, \dots, y_d)$


- The Minkowski metric of order  $r$  is defined by

$$d(x, y) = \sqrt[r]{\sum_i |x_i - y_i|^r}$$

- Common Minkowski metrics:

- ▲ Euclidean( $r=2$ ):  $d(x, y) = \sqrt{\sum_i (x_i - y_i)^2}$ , also called  $L_2$  distance
- ▲ Manhattan distance( $r=1$ ):  $d(x, y) = \sum_i |x_i - y_i|$ , also called  $L_1$  distance
- ▲ “Sup” distance( $r = +\infty$ ):  $d(x, y) = \max_i |x_i - y_i|$ , also called  $L_\infty$  distance

# A simple example



- 1: Euclidean distance:  $\sqrt{4^2 + 3^2} = 5.$
- 2: Manhattan distance:  $4 + 3 = 7.$
- 3: "sup" distance:  $\max\{4, 3\} = 4.$

# Similarities

- Cosine similarity – commonly used to measure document similarity



$$\cos(x, x') = \frac{\langle x \cdot x' \rangle}{|x| \cdot |x'|}$$

- Kernels: RBF (Gaussian) Kernel

$$K(X, X') = \exp \frac{-|X - X'|^2}{2\sigma^2}$$

# Clustering Algorithms


- **Hierarchical algorithms**
  - ▲ Bottom up (agglomerative)
  - ▲ Top down (divisive)
- **Partition algorithms (Flat)**
  - ▲ K-means
  - ▲ Mixture of Gaussians
  - ▲ Spectral clustering





# Hierarchical Clustering

- Given a set of objects, build a tree-based taxonomy



- Hierarchies are a convenient way for organizing information, used frequently by web-portals

# Hierarchical Agglomerative Clustering (HAC)


- ▲ Start with each object in a separate cluster
- ▲ Repeatedly join the closest pair of clusters
- ▲ until there is only one cluster

- The history of merging forms a tree of hierarchy
- **Question**: how to measure the “**closeness**” of two clusters?

# Closest Pair of Clusters?


The distance between two clusters is defined as the distance between:

- Single-link
  - ▲ The nearest pair of points
- Complete-link
  - ▲ The farthest pair of points
- Centroid
  - ▲ The center of gravity
- Average-link
  - ▲ Average of all cross-cluster pairs




# Visualization of the hierarchy: Dendrogram

- Can be used to identify the number of clusters in data
  - ▲ A horizontal cut will create a unique clustering
  - ▲ Moving the cut from root down creates more clusters
  - ▲ Large gaps between the merging nodes indicate a good cutting point




# Dendograms


Single-Link





Complete-Link





# Another example



# Single Link vs. Complete Link



**Single**



**Complete**

- Single-link creates straggly clusters due to chaining effect

# Computational Complexity

- All hierarchical clustering methods need to compute distance of all pairs of  $n$  individual instances which is  $O(n^2)$
- There are  $n-1$  iterations, at each iteration after the merge we must compute the distance between new cluster and all other clusters

$$\sum_{i=2}^{n-1} n - i = O(n^2)$$

- In order to maintain an overall  $O(n^2)$  performance, distance update must be done in constant time – trivial for complete-link and single-link



# Partition Clustering

- Given a data set of  $n$  points, we know that there are  $k$  clusters in the data, how to find these clusters?
- Roughly speaking there are  $O(k^n)$  ways to partition the data, Which one is better?
- One intuition says that we want tight clusters, i.e., points should be in a tight ball
- This leads to the following objective function

$$\sum_{i=1}^k \sum_{x \in C_i} |x - \mu_i|^2 \quad \text{--- squared distance between data point } x \text{ and its cluster center}$$

- Optimizing this objective is a combinatorial optimization problem
  - ▲ Exhaustive search for an optimal solution is not feasible

# Combinatorial optimization: An iterative solution

- ***Initialization:*** Start with a random partition of the data
- ***Iterative step:*** the cluster assignments and cluster centers are updated to improve the objective
- ***Stopping criterion:*** if no improvement can be achieved.

Iterative greedy descent

– convergence is guaranteed, but to local optimal

# K-Means

## Algorithm

**Input** – Desired number of clusters,  $k$

**Initialize** – the  $k$  cluster centers (randomly if necessary)

**Iterate** –


1. Assigning each of the  $N$  data points to its nearest cluster centers
2. Re-estimate the cluster center by assuming that the current assignment is correct

$$\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$$

**Termination** –

If none of the data points changed membership in the last iteration, exit. Otherwise, go to 1

# K-Means Example (K=2)



Pick seeds

Reassign clusters

Compute centroids

Reassign clusters

Compute centroids

Reassign clusters

**Converged!**

# Computational Complexity

- At each iteration:
  - ▲ Reassigning clusters:  $O(kn)$  distance computations
  - ▲ Computing centroids: Each instance vector gets added once to some centroid:  $O(n)$
- Assume these two steps are each done once for  $l$  iterations:  $O(lkn)$ .
- Linear in all relevant factors, assuming a fixed number of iterations, more efficient than  $O(n^2)$  HAC
- Does it always converge?

# K-means Convergence

## Objective

$$\min_{\mu} \min_C \sum_{i=1}^k \sum_{x \in C_i} |x - \mu_i|^2$$

1. Fix  $\mu$ , optimize  $C$ :

***Step 1 of kmeans***

$$\min_C \sum_{i=1}^k \sum_{x \in C_i} |x - \mu_i|^2 = \min_c \sum_i^n |x_i - \mu_{x_i}|^2$$

2. Fix  $C$ , optimize  $\mu$ :

$$\min_{\mu} \sum_{i=1}^k \sum_{x \in C_i} |x - \mu_i|^2$$

- Take partial derivative of  $\mu_i$  and set to zero, we have


$$\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$$

***Step 2 of kmeans***

K-means takes an alternating optimization approach, each step is guaranteed to decrease the objective – thus guaranteed to converge

# Impact of Initial Seeds

- Highly sensitive to the initial seeds



- Multiple random trials: choose the one with best sum of squared loss (important!)
- Heuristics for choosing better centers
  - ▲ choose initial centers to be far apart – furthest first traversal (K-Means++ algorithm)
  - ▲ Initialize with results of other clustering method

# More Comments

- K-Means is exhaustive:
  - ▲ Cluster every data point, no notion of outlier
  - ▲ Outliers cause problems
    - Become singular clusters
    - Bias the centroid estimation
- K-medoids methods is more robust to outliers
  - ▲ Cluster medoid: the point that has minimum sum squared distance to all data points in the cluster
  - ▲ More expensive to compute
    - For each point: sum squared distance with all other pts in cluster  $O(|C|^2)$
- Need to specify  $k$ : difficult in practice
  - ▲ Automatically deciding  $k$ ? more on this later...



# Soft Clustering

- Hard clustering:
  - ▲ Data point is deterministically assigned to one and only one cluster
  - ▲ But in reality clusters may overlap
- Soft-clustering:
  - ▲ Data points are assigned to clusters with certain probabilities
- Model-based clustering

# Aside: Gaussian Bayes Classifier

- We have  $k$  classes in our data
- Each class contains data generated from a particular Gaussian distribution
- Data is generated by
  - ▲ first randomly select one of the classes according to class prior  $p(y)$
  - ▲ draw random samples from the Gaussian distribution of that particular class

$$P(\mathbf{x}, y) = P(\mathbf{x} | y)P(y)$$
$$P(\mathbf{x} | y = i) = \frac{1}{(2\pi)^{d/2} |\Sigma_i|^{1/2}} e^{-\frac{1}{2}(\mathbf{x} - \mu_i)^T \Sigma_i^{-1} (\mathbf{x} - \mu_i)}$$

# Back to Unsupervised Learning

- Now assume we know our data is generated in the same way
- If we know the labels, we can estimate the mean and covariance of the each class using ML (Maximum Likelihood) estimation
  - ▲ Bayes Gaussian Classifier
- But for unsupervised learning, we don't have the labels
- How can we learn the correct model from the incomplete data?

# Gaussian Mixture Model

$$\begin{aligned} P(\mathbf{x}) &= \sum_{i=1}^k P(\mathbf{x}, y = i) \\ &= \sum_{i=1}^k P(\mathbf{x} | y = i) P(y = i) \\ &= \sum_{i=1}^k \alpha_i P(\mathbf{x} | \theta_i) \end{aligned}$$

$\alpha_i = p(y=i)$ : the class prior  
Mixing parameter

$\theta_i = \{\mu_i, \Sigma_i\}$

Goal of unsupervised learning:

- Given a set of  $\mathbf{x}$ 's, estimate  $\{\alpha_1, \dots, \alpha_k, \theta_1, \dots, \theta_k\}$
- Once the model is identified, we can compute  $p(y = i | \mathbf{x})$  for  $i = 1, \dots, k$

# Maximum Marginal Likelihood

$$\begin{aligned}\arg \max_{\theta} \prod_j P(\mathbf{x}^j) &= \arg \max_{\theta} \prod_j \sum_{i=1}^k P(\mathbf{x}^j, y^j = i) \\ &= \arg \max_{\theta} \sum_{j=1}^n \underbrace{\log \sum_{i=1}^k P(\mathbf{x}^j, y^j = i)}\end{aligned}$$

log sum is difficult to optimize !

Gradient ascent is doable but very inefficient

# Expectation Maximization (EM)

- A highly used approach for dealing with hidden (missing) data
  - ▲ Here the cluster labels are hidden
- Much simpler than gradient methods
- It is an iterative algorithm that starts with some initial guess of the model parameters
- Iteratively performs two linked steps:
  - ▲ **Expectation (E-step)**: given current model parameters  $\lambda_t$ , compute the expectation for the hidden (missing) data
  - ▲ **Maximization (M-step)**: re-estimate the parameters  $\lambda_{t+1}$  assuming that the expected values computed in the E-step are the true values
- We will first show how it works for mixture of Gaussian

# EM – simple case

- A simple case:
  - ▶ We have unlabeled data  $x^1, \dots, x^m$
  - ▶ We know there are  $k$  classes
  - ▶ We know  $\alpha_1 = P(y = 1), \dots, \alpha_k = P(y = k)$
  - ▶ We don't know  $\mu_1 \dots \mu_k$ , but know the common variance  $\sigma^2$

Start with an initial guess for  $\mu_1, \dots, \mu_k$ ,

1. If we know  $\mu_1, \dots, \mu_k$ , we can easily compute probability that a point  $x^j$  belongs to class  $i$ :

$$p(y = i | x^j) \propto \exp\left(-\frac{1}{2\sigma^2} |x^j - \mu_i|^2\right) p(y = i)$$

Simply evaluate this, then normalize

**E-step**

2. If we know *the* probability that each point belongs to each class, we can estimate the  $\mu_1, \dots, \mu_k$

$$\mu_i = \frac{\sum_{j=1}^m p(y=i|x^j) x^j}{\sum_{j=1}^m p(y=i|x^j)}$$

**M-step**

# EM – Axis-aligned Gaussian

- We have unlabeled data  $x^1, \dots, x^m$
- We know there are  $k$  classes
- We know that the Gaussians are axis aligned

$$\Sigma_i = \begin{pmatrix} \sigma_{i,1}^2 & 0 & 0 & \dots & 0 & 0 \\ 0 & \sigma_{i,2}^2 & 0 & \dots & 0 & 0 \\ 0 & 0 & \sigma_{i,3}^2 & \dots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \dots & \sigma_{i,m-1}^2 & 0 \\ 0 & 0 & 0 & \dots & 0 & \sigma_{i,m}^2 \end{pmatrix}$$

Start with an initial guess for  $\mu_1, \dots, \mu_k, \Sigma_1, \dots, \Sigma_k, \alpha_1, \dots, \alpha_k$ ,

1. If we know the parameters, we can easily compute probability that a point  $x^j$  belongs to class  $i$ :

$$p(y = i | x^j) \propto p(x^j | \mu_i, \Sigma_i) p(y = i)$$

Simply evaluate this, then normalize

**E-step**

2. If we know *the* probability that each point belongs to each class, we can estimate the  $\mu_1, \dots, \mu_k, \Sigma_1, \dots, \Sigma_k, \alpha_1, \dots, \alpha_k$ ,

$$\mu_i = \frac{\sum_{j=1}^m p(y=i|x^j) x^j}{\sum_{j=1}^m p(y=i|x^j)} \quad \alpha_i = \frac{\sum_{j=1}^m p(y=i|x^j)}{m} \quad \sigma_{il}^2 = \frac{\sum_{j=1}^m p(y=i|x^j) (x_l^j - \mu_{il})^2}{\sum_{j=1}^m p(y=i|x^j)}$$

**M-step**



# EM – General Gaussian

Start with an initial guess for  $\mu_1, \dots, \mu_k, \Sigma_1, \dots, \Sigma_k, \alpha_1, \dots, \alpha_k$ ,

1. If we know the parameters, we can easily compute probability that a point  $x^j$  belongs to class  $i$ :

$$p(y = i | x^j) \propto p(x^j | \mu_i, \Sigma_i) p(y = i)$$

Simply evaluate this, then normalize

**E-step**


2. If we know *the* probability that each point belongs to each class, we can estimate the  $\mu_1, \dots, \mu_k, \Sigma_1, \dots, \Sigma_k, \alpha_1, \dots, \alpha_k$ ,

$$\mu_i = \frac{\sum_{j=1}^m p(y=i|x^j) x^j}{\sum_{j=1}^m p(y=i|x^j)} \quad \alpha_i = \frac{\sum_{j=1}^m p(y=i|x^j)}{m}$$


$$\Sigma_i = \frac{\sum_{j=1}^m p(y=i|x^j) (x^j - \mu_i)(x^j - \mu_i)^T}{\sum_{j=1}^m p(y=i|x^j)}$$

**M-step**


# Gaussian Mixture Example: Start




# After first iteration




# After 2nd iteration




# After 3rd iteration




# After 4th iteration



# After 5th iteration




# After 6th iteration





# After 20th iteration



# Behavior of EM

- It is guaranteed to converge
  - ▲ Convergence proof is based on the fact that  $P(x|\theta)$  must increase or remain same between iterations (not obvious)
  - ▲ In practice it may converge slowly, one can stop early if the change in log-likelihood is smaller than a threshold
- It converges to a local optimum
  - ▲ Multiple restart is recommended